

Problem Set 1

Rajvi Jasani

GitHub Repository

This is the link to my GitHub repository <https://github.com/rajvijasani/STATS506-Problem-Set-1.git>

Problem 1 - Wine Data

```
setwd("E:/UM/STATS 506/Repos/STATS506-Problem-Set-1")

# Assigning column names for the data frame
colNames <- c(
  "wineClass",
  "alcohol",
  "malicAcid",
  "ash",
  "alcalinityOfAsh",
  "magnesium",
  "totalPhenols",
  "flavanoids",
  "nonflavanoidPhenols",
  "proanthocyanins",
  "colorIntensity",
  "hue",
  "odOfDilutedWines",
  "proline"
)

# Importing the data set
wine <- read.table("data/wine.data", sep = ",", col.names = colNames)
```

```
# Checking number of data samples for each class of wine
classCount <- table(wine$wineClass)
print(classCount)
```

```
1 2 3
59 71 48
```

The number of data samples for each class of wine are the same as mentioned in the wine.names file.

```
# Calculating correlation of alcohol content and color intensity
corRel <- cor(wine$alcohol, wine$colorIntensity)
print(corRel)
```

```
[1] 0.5463642
```

The correlation coefficient (r) between alcohol content and color intensity of wine is 0.5463642. As r is positive, we can say that there is a positive correlation, i.e. if one variable increases, the other also increases.

```
# Splitting the data frame based on wine classes
# This returns a list of 3 lists grouped on the 3 wine classes
splitWine <- split(wine, wine$wineClass)
```

Attribution of Source: Asked UM GPT to suggest ways to split dataset by groups to enable me to apply correlation function to each.

```
# Using lapply to calculate correlation between alcohol and color intensity for each list generated
#' Anonymous function to calculate correlation
#' @param splitWineList list (subset of the original dataset) containing samples of same class
#' @return correlation coefficient between alcohol and color intensity from data in the given list
corRelByClass <- lapply(splitWine, function(splitWineList) {
  return(cor(splitWineList$alcohol, splitWineList$colorIntensity))
})
print(corRelByClass)
```

```
$`1`  
[1] 0.4082913
```

```
$`2`  
[1] 0.2697891
```

```
$`3`  
[1] 0.3503777
```

From the output, we can observe that class 1 of wine has the highest correlation while class 2 has the lowest correlation between alcohol content and color intensity.

```
# Finding the index of row with highest value of color intensity  
maxColorIntensityIndex <- which.max(wine$colorIntensity)  
# Using the index to get corresponding alcohol content value  
print(wine$alcohol[maxColorIntensityIndex])
```

```
[1] 14.34
```

The alcohol content of the sample of wine with the highest color index is 14.34.

```
# Counting the number of samples where proanthocyanins > ash  
counter = 0  
for (i in 1:nrow(wine)) {  
  if (wine$proanthocyanins[i] > wine$ash[i]) {  
    counter <- counter + 1  
  }  
}  
# Calculating the percentage  
percentageWine <- (counter / nrow(wine)) * 100  
print(percentageWine)
```

```
[1] 8.426966
```

Therefore, approximately 8.43% of the samples have more proanthocyanins than ash.

```
# Calculating average of each variable except wineClass  
overallAverage <- colMeans(wine[, -1])  
# Converting above obtained vector into a dataframe  
overallAverageDf <- as.data.frame(t(overallAverage), row.names = "Overall")
```

```
# Calculating average of each variable grouped by wineClass
groupAverage <- aggregate(. ~ wineClass, wine, mean)
# Removing the wineClass column from the data frame
groupAverage <- groupAverage[, -1]
# Binding row-wise, the two dataframes obtained above
averagesTable <- rbind(overallAverageDf, groupAverage)
```

Attribution of Source: Asked UM GPT to suggest ways to get average of all variables by group and for ways to combine and display the table of averages.

```
print(averagesTable)
```

	alcohol	malicAcid	ash	alcalinityOfAsh	magnesium	totalPhenols
Overall	13.00062	2.336348	2.366517	19.49494	99.74157	2.295112
1	13.74475	2.010678	2.455593	17.03729	106.33898	2.840169
2	12.27873	1.932676	2.244789	20.23803	94.54930	2.258873
3	13.15375	3.333750	2.437083	21.41667	99.31250	1.678750

	flavanoids	nonflavanoidPhenols	proanthocyanins	colorIntensity	hue
Overall	2.0292697	0.3618539	1.590899	5.058090	0.9574494
1	2.9823729	0.2900000	1.899322	5.528305	1.0620339
2	2.0808451	0.3636620	1.630282	3.086620	1.0562817
3	0.7814583	0.4475000	1.153542	7.396250	0.6827083

	odOfDilutedWines	proline
Overall	2.611685	746.8933
1	3.157797	1115.7119
2	2.785352	519.5070
3	1.683542	629.8958

The table above lists the overall as well group-wise averages for each variable.

```
# Using the splitWine list created above to extract values of total phenol for each class
# Pairwise t-test on class 1 and class 2
test12 <- t.test(splitWine[[1]]$totalPhenols, splitWine[[2]]$totalPhenols)
print(test12)
```

Welch Two Sample t-test

```
data: splitWine[[1]]$totalPhenols and splitWine[[2]]$totalPhenols
t = 7.4206, df = 119.14, p-value = 1.889e-11
```

```
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.4261870 0.7364055
sample estimates:
mean of x mean of y
 2.840169  2.258873
```

```
# Pairwise t-test on class 1 and class 3
test13 <- t.test(splitWine[[1]]$totalPhenols, splitWine[[3]]$totalPhenols)
print(test13)
```

Welch Two Sample t-test

```
data: splitWine[[1]]$totalPhenols and splitWine[[3]]$totalPhenols
t = 17.12, df = 98.356, p-value < 2.2e-16
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 1.026801 1.296038
sample estimates:
mean of x mean of y
 2.840169  1.678750
```

```
# Pairwise t-test on class 2 and class 3
test23 <- t.test(splitWine[[2]]$totalPhenols, splitWine[[3]]$totalPhenols)
print(test23)
```

Welch Two Sample t-test

```
data: splitWine[[2]]$totalPhenols and splitWine[[3]]$totalPhenols
t = 7.0125, df = 116.91, p-value = 1.622e-10
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.4162855 0.7439610
sample estimates:
mean of x mean of y
 2.258873  1.678750
```

Attribution of Source: Asked UM GPT for inbuilt function to calculate t-test.

Problem 2 - AskAManager.org Data

```
# Importing data
salary = read.csv("data/AskAManager.csv")
# Re-assigning names to columns
names(salary) <- c(
  "index",
  "timestamp",
  "age",
  "industry",
  "jobTitle",
  "jobDescription",
  "annualSalary",
  "bonus",
  "currency",
  "otherCurrency",
  "additionalDescription",
  "country",
  "state",
  "city",
  "overallExperience",
  "fieldExperience",
  "levelOfEducation",
  "gender",
  "race"
)
print(names(salary))
```

```
[1] "index"           "timestamp"       "age"
[4] "industry"        "jobTitle"        "jobDescription"
[7] "annualSalary"    "bonus"           "currency"
[10] "otherCurrency"   "additionalDescription" "country"
[13] "state"           "city"            "overallExperience"
[16] "fieldExperience" "levelOfEducation" "gender"
[19] "race"
```

Number of samples in original dataset

```
print(nrow(salary))
```

```
[1] 28062
```

```
# Choosing samples that have salary in USD
salaryRestricted <- salary[salary$currency == "USD",]
```

Number of samples in restricted dataset (less than the original)

```
print(nrow(salaryRestricted))
```

```
[1] 23374
```

Cleaning based on age and experience

```
# Need to look at all the possible values for age, overall experience and field experience
print(table(salaryRestricted$age))
```

18-24	25-34	35-44	45-54	55-64	65 or over	under 18
991	10386	8323	2700	875	86	13

```
print(table(salaryRestricted$overallExperience))
```

1 year or less	11 - 20 years	2 - 4 years	21 - 30 years
410	8137	2450	3066
31 - 40 years	41 years or more	5-7 years	8 - 10 years
735	112	4016	4448

```
print(table(salaryRestricted$fieldExperience))
```

1 year or less	11 - 20 years	2 - 4 years	21 - 30 years
1180	5523	5045	1573
31 - 40 years	41 years or more	5-7 years	8 - 10 years
333	34	5460	4226

Mapping range values to numeric values for age, overall experience and field experience

```

#' Function to map age values given in range to minimum limit of the range in numeric format
#'
#' @param ageRanges takes column of age values from the dataset
#'
#' @return column of mapped/factored age values in numeric format
mapAge <- function(ageRanges) {
  ageFactors <- c(
    "under 18" = 0,
    "18-24" = 18,
    "25-34" = 25,
    "35-44" = 35,
    "45-54" = 45,
    "55-64" = 55,
    "65 or over" = 65
  )
  return(as.numeric(ageFactors[ageRanges]))
}

#' Function to map experience values given in range to maximum limit of the range in numeric
#'
#' @param experienceRanges takes column of overall or field experience values from the dataset
#'
#' @return column of mapped/factored experience values in numeric format
mapExperience <- function(experienceRanges) {
  experienceFactors <- c(
    "1 year or less" = 1,
    "2 - 4 years" = 4,
    "5-7 years" = 7,
    "8 - 10 years" = 10,
    "11 - 20 years" = 20,
    "21 - 30 years" = 30,
    "31 - 40 years" = 40,
    "41 years or more" = 41
  )
  return(as.numeric(experienceFactors[experienceRanges]))
}

```

Attribution of Source: Asked UM GPT to suggest ways to factor character ranges to numeric values.

Mapping/Factoring columns


```
salaryRestricted$minAge <- sapply(salaryRestricted$age, mapAge)
salaryRestricted$maxOverallExperience <- sapply(salaryRestricted$overallExperience, mapExperience)
salaryRestricted$maxFieldExperience <- sapply(salaryRestricted$fieldExperience, mapExperience)
```

Cleaning data for impossible age and experience combinations

```
salaryCleaned <- salaryRestricted[salaryRestricted$minAge >= 18 &
                                   salaryRestricted$maxOverallExperience <= (salaryRestricted$maxOverallExperience - 1) &
                                   salaryRestricted$maxFieldExperience <= (salaryRestricted$maxFieldExperience - 1), ]
```

Number of samples in cleaned dataset (less than the restricted)

```
print(nrow(salaryCleaned))
```

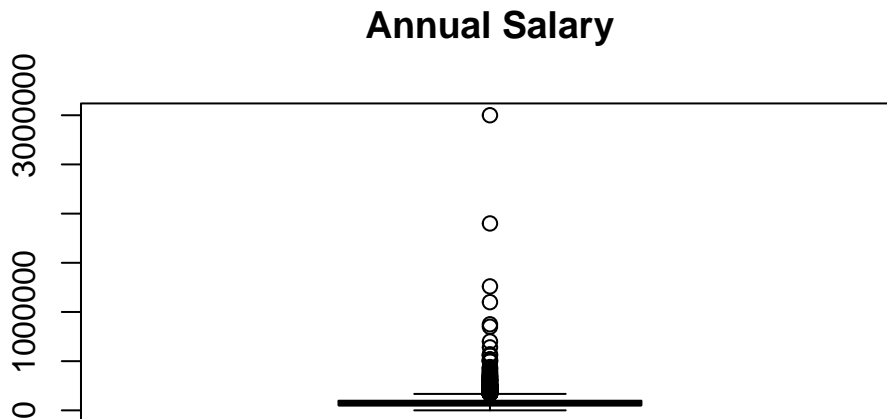
```
[1] 7834
```

Visualizing salaries to understand the trends

```
summary(salaryCleaned$annualSalary)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0	52000	70000	82387	98000	3000000

```
boxplot(salaryCleaned$annualSalary, main = "Annual Salary")
```



From summary and box plot of the annual salaries, we can identify some extreme values. Salaries equal to 0 USD might be a mistake or we might have to consider what the person has written in additional description field. Salaries as high as 3000000 USD are possible but rare. Thus it would be safe to restrict lower limit to a little less than the 1st Quartile value, say 30,000 USD and the upper limit to a bit more than the 3rd Quartile value, say 300,000 USD. This assumption is based on summary statistics, plots and my personal intuition. To further justify it or modify it, we can look into sources such as US Bureau of Labor Statistics <https://www.bls.gov/>. However I have not chosen to do that as the dataset covers jobs from many different industries and many different positions which makes the task difficult.

```
minThreshold <- 30000
maxThreshold <- 300000
salaryFinal <- salaryCleaned[salaryCleaned$annualSalary >= minThreshold &
                             salaryCleaned$annualSalary <= maxThreshold, ]
```

Number of samples in final dataset (less than the cleaned)

```
print(nrow(salaryFinal))
```

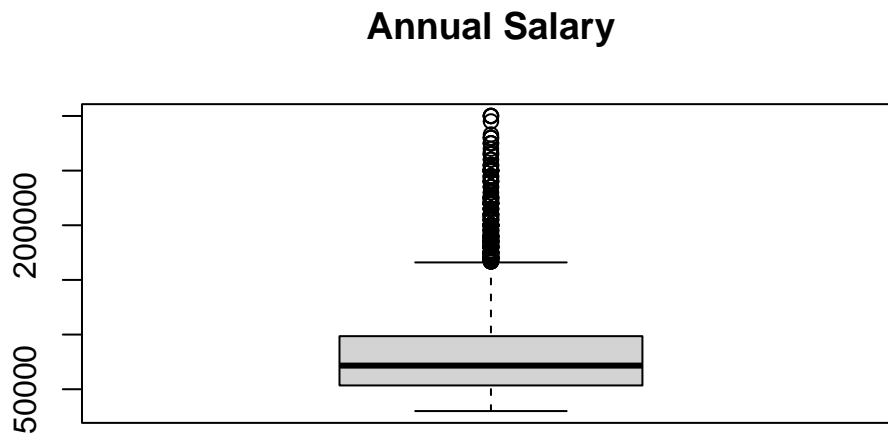
```
[1] 7573
```

Final summary statistics and box plot

```
summary(salaryFinal$annualSalary)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
30000	53461	71600	81218	98530	300000

```
boxplot(salaryFinal$annualSalary, main = "Annual Salary")
```



Problem 3 - Palindromic Numbers

```
#' Function to check if given input is Palindromic and output the reversed number
#'
#' @param num positive integer input
#'
#' @return named list with 2 elements, a logical "isPalindromic" indicating if input is a Pal
isPalindromic <- function(num) {
  # Checking if input is numeric; if not, converting it to numeric
  if (!is.numeric(num)) {
    num <- as.numeric(num)
  }
}
```

```

# Checking if input is positive; if not, showing relevant error
if (num < 0) {
  stop("number must be positive")
}
# Checking if input is an integer; if not, showing relevant error
if (num != floor(num)) {
  stop("number must be an integer")
}
result <- list()
numOG <- num
numReverse <- 0

# Mathematical way of reversing a number
while (num > 0) {
  # Extracting last digit
  remainder <- num %% 10
  # Appending it to reverse of the number
  numReverse <- numReverse * 10 + remainder
  # Discarding the last digit in the original number
  num <- num %/% 10
}
# To handle cases with trailing zeros because
# I have used mathematical approach of reversing a number
if (nchar(as.character(numOG))!=nchar(as.character(numReverse))) {
  result["warning"] <- "Your input has trailing zeros. The computed reverse might not be a"
}
if (numReverse == numOG) {
  result["isPalindromic"] <- TRUE
  result["reversed"] <- numReverse
}
else{
  result["isPalindromic"] <- FALSE
  result["reversed"] <- numReverse
}
return(result)
}

```

Attribution of Source: Asked UM GPT for better way to check if a number is an integer instead of using `is.integer()` function.

```

#' Function to get the next palindromic number strictly greater than the input
#'

```

```

#' @param num positive integer input
#'
#' @return vector containing the next palindromic number
nextPalindrome <- function(num) {
  # Checking if input is numeric; if not, converting it to numeric
  if (!is.numeric(num)) {
    num <- as.numeric(num)
  }
  # Checking if input is positive; if not, showing relevant error
  if (num < 0) {
    stop("number must be positive")
  }
  # Checking if input is an integer; if not, showing relevant error
  if (num != floor(num)) {
    stop("number must be an integer")
  }
  # We need a number strictly greater than the input.
  # In some cases the input itself might be a palindrome.
  # To avoid that, we start checking with the number after the input
  num <- num + 1
  # While the number is not a palindrome,
  while (isPalindromic(num)["isPalindromic"] == FALSE) {
    # Go to the next number
    num <- num + 1
  }
  return (c(num))
}

```

Finding the next palindrome for each of the following

```
nextPalindrome(391)
```

```
[1] 393
```

```
nextPalindrome(9928)
```

```
[1] 9999
```

```
nextPalindrome(19272719)
```

```
[1] 19277291
```

```
nextPalindrome(109)
```

```
[1] 111
```

```
nextPalindrome(2)
```

```
[1] 3
```

Attribution of Source: Asked ChatGPT if a single digit number can be called a palindrome. It can be.